

Distributed programming

jan-feb 2026, subject no. 3

1. (3p) Consider the following excerpt from a program that is supposed to merge-sort a vector. The function `worker()` is called in all processes except process 0, the function `mergeSort()` is called from process 0 (and from the places described in this excerpt), the function `mergeSortLocal()` sorts the specified vector and the function `mergeParts()` merges two sorted adjacent vectors, given the pointer to the first element, the total length and the length of the first vector.

```

1 void mergeSort(int* v, int dataSize, int myId, int nrProc) {
2     if(dataSize <= 1) {
3         mergeSortLocal(v, dataSize);
4     } else {
5         int halfProc = (nrProc+1) / 2; → how many procs the "Child" gets
6         int child = myId+halfProc; → rankId of the leader of the child group
7         int halfSize = (dataSize+1) / 2;
8         MPI_Ssend(halfProc, 1, MPI_INT, child, 1, MPI_COMM_WORLD);
9         MPI_Ssend(&halfSize, 1, MPI_INT, child, 2, MPI_COMM_WORLD);
10        MPI_Ssend(v, halfSize, MPI_INT, child, 3, MPI_COMM_WORLD);
11        mergeSort(v+halfSize, dataSize/2, myId, halfProc);
12        MPI_Recv(v, halfSize, MPI_INT, child, 4, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
13        mergeParts(v, dataSize, halfSize);
14    }
15 }
16 void worker(int myId) {
17     MPI_Status status;
18     int dataSize, nrProc;
19     → MPI_Recv(&nrProc, 1, MPI_INT, MPI_ANY_SOURCE, 1, MPI_COMM_WORLD, &status);
20     auto parent = status.MPI_SOURCE;
21     → MPI_Recv(&dataSize, 1, MPI_INT, parent, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
22     std::vector v(dataSize);
23     → MPI_Recv(v.data(), dataSize, MPI_INT, parent, 3, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
24     mergeSort(v.data(), dataSize, myId, nrProc); ←
25     MPI_Ssend(v.data(), dataSize, MPI_INT, parent, 4, MPI_COMM_WORLD);
26 }
```

Which of the following issues are present? Describe the changes needed to solve them.

- A: some worker processes are not used if the number of processes is not a power of 2.
- B: the application can deadlock if the length of the vector is smaller than the number of MPI processes.
- C: the application can deadlock if the number of processes is not a power of 2. ✓
- D: the application can produce a wrong result if the input vector size is not a power of 2.

$$v = [5, 3]$$

$$nrProc = 6$$

$$ms([6, 3], 4, 0, 4)$$

$$\begin{matrix} V & P \\ 1 & 3 \end{matrix}$$

feb 2022, subject no. 1

1. (3p) Consider the following excerpt from a program that is supposed to merge-sort a vector. The function `worker()` is called in all processes except process 0, the function `mergeSort()` is called from process 0 (and from the places described in this excerpt), the function `mergeSortLocal()` sorts the specified vector and the function `mergeParts()` merges two sorted adjacent vectors, given the pointer to the first element, the total length and the length of the first vector.

```

1 void mergeSort(int* v, int dataSize, int myId, int nrProc) {
2     if(nrProc == 1 || dataSize <= 1) {
3         mergeSortLocal(v, dataSize);
4     } else {
5         int halfLen = dataSize / 2;
6         int halfProc = nrProc / 2;
7         int child = myId+halfProc;
8         MPI_Ssend(&halfLen, 1, MPI_INT, child, 1, MPI_COMM_WORLD);
9         MPI_Ssend(&halfProc, 1, MPI_INT, child, 2, MPI_COMM_WORLD);
10        MPI_Ssend(v, halfSize, MPI_INT, child, 3, MPI_COMM_WORLD);
11        mergeSort(v+halfSize, dataSize-halfSize, myId, nrProc);
12        MPI_Recv(v, halfSize, MPI_INT, child, 4, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
13        mergeParts(v, dataSize, halfSize);
14    }
15 }
16 void worker(int myId) {
17     MPI_Status status;
18     int dataSize, nrProc;
19     MPI_Recv(&dataSize, 1, MPI_INT, MPI_ANY_SOURCE, 1, MPI_COMM_WORLD, &status);
20     auto parent = status.MPI_SOURCE;
21     MPI_Recv(&nrProc, 1, MPI_INT, parent, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
22     std::vector v(dataSize);
23     MPI_Recv(v.data(), dataSize, MPI_INT, parent, 3, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
24     mergeSort(v.data(), dataSize, myId, nrProc);
25     MPI_Ssend(v.data(), dataSize, MPI_INT, parent, 4, MPI_COMM_WORLD);
26 }
```

$$\begin{matrix} V & P \\ 1 & 3 \end{matrix}$$

Which of the following issues are present? Describe the changes needed to solve them.

- A: the application can deadlock if the length of the vector is smaller than the number of MPI processes.
- B: the application can produce a wrong result if the input vector size is not a power of 2.
- C: some worker processes are not used if the number of processes is not a power of 2.
- D: the application can deadlock if the number of processes is not a power of 2.

2. (3p) Consider the following code for implementing a future mechanism (the `set()` function is guaranteed to be called exactly once by the user code)

```

1 template<typename T>
2 class Future {
3     T val;
4     bool hasValue;
5     mutex mtx;
6     condition_variable cv;
7 public:
8     Future() :hasValue(false) {}
9     void set(T v) {
10
11         cv.notify_all();
12         unique_lock<mutex> lck(mtx);
13         hasValue = true;
14         val = v;
15     }
16     T get() {
17         unique_lock<mutex> lck(mtx);
18         while(!hasValue) {
19             cv.wait(lck);
20         }
21         return value;
22     }
23 };

```

Which of the following are true? Give a short explanation.

- A: [issue] a call to `get()` can deadlock if simultaneous with the call to `set()` \top
- B: [issue] a call to `get()` can deadlock if called after `set()` \top
- C: [issue] a call to `get()` can return an uninitialized value if simultaneous with the call to `set()` \top
- D: [issue] simultaneous calls to `get()` and `set()` can make future calls to `get()` deadlock
- E: [issue] a call to `get()` can deadlock if called before `set()`
- F: [fix] a possible fix is to remove the line 11
- G: [fix] a possible fix is to interchange lines 12 and 13
- H: [fix] a possible fix is to reorder lines 10–13 in the order 11, 13, 12, 10
- I: [fix] a possible fix is to interchange lines 10 and 11
- J: [fix] a possible fix is to unlock the mutex just before line 18 and to lock it back just afterwards

Exam to Parallel and Distributed Programming ian-feb 2025, subject no. 2

1. (3p) Consider the following excerpt from a program that is supposed to compute the scalar product of two vectors of the same length. The function `worker()` is called in all processes except process 0, the function `product()` is called from process 0.

```

P6
1 int product(int nrProc, std::vector<int> const& p, std::vector<int> const& q) {
2     int chunkSize = (p.size() + nrProc - 1) / nrProc; • 1
3     std::vector<int> partResults(nrProc);
4     partProd(chunkSize, p.data(), q.data(), partResults.data());
5     int sum = 0;
6     for(int const& v : partResults) sum += v;
7     return sum;
8 }
9 void worker(int myId, int nrProc) {
10     partProd(0, nullptr, nullptr, nullptr);
11 }
12 void partProd(int chunkSize, int const* p, int const* q, int* r) {
13     MPI_Bcast(&chunkSize, 1, MPI_INT, 0, MPI_COMM_WORLD);
14     std::vector<int> pp(chunkSize);
15     std::vector<int> qq(chunkSize);
16     MPI_Scatter(p, chunkSize, MPI_INT, pp.data(), chunkSize, MPI_INT, 0, MPI_COMM_WORLD);
17     MPI_Scatter(q, chunkSize, MPI_INT, qq.data(), chunkSize, MPI_INT, 0, MPI_COMM_WORLD);
18     int sum = 0;
19     for(int i=0 ; i<chunkSize ; ++i) {
20         sum += pp[i]*qq[i];
21     }
22     MPI_Gather(&sum, 1, MPI_INT, r, 1, MPI_INT, 0, MPI_COMM_WORLD);
23 }

```

Which of the following issues are present if the vector lengths are not a multiple of the number of MPI processes? Describe the changes needed to solve them.

$$P = \overbrace{[10, 20, 30, 40]}$$

$\text{proc} = 3$

$$(4+3-1)/3$$

≈ 2

2 2 2

psize
1
mproc
10

- A: the application can attempt an illegal memory access.
 B: the application can deadlock. False
 C: some worker processes are not used. False
 D: some terms are added twice.
 E: some terms are not added at all.
 F: the scalar product is incorrectly computed in some other way.

2. (3p) Consider the following code for enqueueing a continuation on a future (the `set()` function is guaranteed to be called exactly once by the user code):

```

1  template<typename T>
2  class Future {
3      list<function<void(T)>> continuations;
4      T val;
5      bool hasValue;
6      mutex mtx;
7  public:
8      Future() :hasValue(false) {}
9
10     void set(T v) {
11         unique_lock<mutex> lck(mtx);
12         hasValue = true;
13         val = v;
14         lck.unlock();
15         for(function<void(T)>& f : continuations) {
16             f(v);
17         }
18         continuations.clear();
19     }
20     void addContinuation(function<void(T)> f) {
21         if(hasValue) {
22             f(val);
23         } else {
24             unique_lock<mutex> lck(mtx);
25             continuations.push_back(f);
26         }
27     }
28 };
  
```

Which of the following are true? Give a short explanation.

- A: [issue] a call to `set()` can deadlock if simultaneous with the call to `addContinuation()`;
 B: [issue] two simultaneous calls to `addContinuation()` may deadlock;
 C: [issue] simultaneous calls to `addContinuation()` and `set()` may lead to continuations that are executed twice;
 D: [fix] two simultaneous calls to `addContinuation()` may lead to a corrupted `continuations` vector;
 E: [issue] simultaneous calls to `addContinuation()` and `set()` may lead to continuations that are never executed;
 F: [fix] a possible fix is to move the content of line 13 to between lines 17 and 18;
 G: [fix] a possible fix is to move the content of line 23 to between lines 19 and 20;
 H: [fix] a possible fix is to move the content of line 23 to between lines 19 and 20 and add an `unlock()` call between lines 20 and 21;
 I: [fix] a possible fix is to move the content of line 13 to between lines 17 and 18 and to move line 23 between lines 19 and 20;
 J: [fix] a possible fix is to move the content of lines 11 and 13 (in this order) to between lines 17 and 18;

The following code computes the product of 2 matrices. The program

2. (3p) Consider the following code for a queue with multiple producers and consumers. The `push()` function is guaranteed to be called exactly once by the user code, and on

The `close()` function is guaranteed to be called exactly once by the user code, and `enqueue()` will not be called after that. `dequeue()` is supposed to block if the queue is empty and to return an empty optional if the queue is closed and all the elements have been dequeued.

```
1 template<typename T>
2 class ProducerConsumerQueue {
3     list<T> items;
4     bool isClosed = false;
5     condition_variable cv;
6
7     mutex mtx;
8
9     public:
10    void enqueue(T v) {
11        unique_lock<mutex> lck(mtx);
12        items.push_back(v);
13        cv.notify_one();
14    }
15    optional<T> dequeue() {
16        unique_lock<mutex> lck(mtx);
17        while(items.empty() && !isClosed) {
18            cv.wait(lck);
19        }
20        lck.unlock();
21        if(!items.empty()) {
22            optional<T> ret(items.front());
23            items.pop_front();
24            return ret;
25        }
26        return optional<T>();
27    }
28    void close() {
29        unique_lock<mutex> lck(mtx);
30        isClosed = true;
31        cv.notify_all();
32    }
33};
```

Which of the following are true? Give a short explanation.

- A: [issue] a call to dequeue() can result data corruption or undefined behavior if simultaneous with the call to enqueue();
 - X [issue] a call to dequeue() can deadlock if simultaneous with the call to enqueue();
 - X [issue] two simultaneous calls to enqueue() may deadlock;
 - D: [issue] two simultaneous calls to dequeue() may result in corrupted items list;
 - X [issue] two simultaneous calls to dequeue() may deadlock;
 - X [fix] move the content of line 14 to between lines 15 and 16 and line 18 between 16 and 17
 - X [fix] eliminate lines 14 and 18
 - H: [fix] remove line 18 and insert copies of it between lines 21 and 22 and between 23 and 24
 - X [fix] unlock the mutex before line 16 and lock it back after line 16
 - X [fix] apply fix H above and, additionally, move line 14 in place of line 18